

CT421: Artificial Intelligence

Project 1: Genetic Algorithm for Exam Timetabling

January 28, 2026

1 Overview

As introduced in class, you will implement a genetic algorithm to solve the exam timetabling problem. This is a well known constraint satisfaction and optimisation problem with many real-world applications.

1.1 Learning Objectives

- Implement a genetic algorithm
- Design appropriate chromosome representations for scheduling problems
- Handle hard and soft constraints in optimisation
- Analyze algorithm performance and parameter sensitivity
- Explore the ability of the algorithm to maintain diversity in the search.

2 Problem Definition

2.1 Input

Given:

- N exams to be scheduled
- K available time slots ($1 \dots K$)
- M students
- An enrollment matrix E where $E[i][j] = 1$ if student i is enrolled in exam j , and 0 otherwise

Note that by varying K , M , N and the values in the enrollment matrix we can vary the difficulty of the problem.

2.2 Objective

Find an assignment of exams to time slots that satisfies all hard constraints and minimizes soft constraint violations.

2.3 Constraints

Hard Constraint (must be satisfied):

- No student may have two exams scheduled in the same time slot

Soft Constraint (minimize): Examples:

- Total number of consecutive exams for students should be minimised.

3 Input File Format

Test instances are provided in plain text format:

```
N K M
e_00 e_01 e_02 ... e_0(N-1)
e_10 e_11 e_12 ... e_1(N-1)
...
e_(M-1)0 e_(M-1)1 ... e_(M-1)(N-1)
```

Where:

- First line: N (number of exams), K (number of slots), M (number of students)
- Following M lines: enrollment matrix (1 if student takes exam, 0 otherwise)

3.1 Example Instance: “tinyexample.txt”

```
4 3 5
1 1 0 0
0 1 1 0
0 0 1 1
1 0 0 1
0 1 0 1
```

This represents:

- 4 exams, 3 time slots, 5 students
- Student 0: enrolled in exams $\{0, 1\}$
- Student 1: enrolled in exams $\{1, 2\}$
- Student 2: enrolled in exams $\{2, 3\}$
- Student 3: enrolled in exams $\{0, 3\}$
- Student 4: enrolled in exams $\{1, 3\}$

4 Implementation Requirements

You must implement the following components:

4.1 Suggested Functions

1. `read_instance(filename)` – Parse input file
2. `initialize_population(pop_size)` – Create random initial solutions
3. `evaluate_fitness(solution)` – Calculate fitness value
4. `select_parents(population)` – Choose parents for reproduction
5. `crossover(parent1, parent2)` – Create offspring from two parents
6. `mutate(solution)` – Apply random changes to a solution
7. `run_ga()` – Main genetic algorithm loop

4.2 Genetic Algorithm Parameters

You should experiment with different parameter values. Suggested starting values:

- Population size: 100
- Number of generations: 500
- Crossover rate: 0.8
- Mutation rate: 0.05
- Selection method: Tournament selection (size 3)

5 Output Requirements

Your program should output:

1. The best solution found (array of time slot assignments)
2. The fitness of the best solution
3. Number of hard constraint violations (should be 0 for valid solutions)
4. Measure of soft constraints violated
5. A plot showing fitness over generations

6 Tips and Suggestions

- Start with the small instances to debug your implementation
- Plot fitness over generations to ensure your GA is improving
- If your GA gets stuck, try increasing mutation rate or population size
- Consider implementing elitism (keeping the best solution across generations)
- I will upload some test cases later.